



CANCER
RESEARCH
UK

CAMBRIDGE
INSTITUTE

CRUK cluster practical sessions

Part I – processes & scripts



UNIVERSITY OF
CAMBRIDGE

login

Log in to the head node, uk-cri-lcst01, using **ssh** and your usual user name & password.

```
CRUK Cambridge Institute - HPC Cluster

Disk quotas for group cr-isd-sot (gid 659):
  Filesystem kbytes  quota  limit  grace  files  quota  limit  grace
    /lustre 2204120      0      0      -     92      0      0      -

Disk quotas for user user (uid 10316):
  Filesystem blocks  quota  limit  grace  files  quota  limit  grace
/dev/mapper/criClusterHome01-Data01
      2098824 10485760 12582912      82      0      0

Disk quotas for group cr-isd-sot (gid 659): none

[user@cluster ~]$
```

You're ready to start.

navigate

Find out where you are using **pwd**.

Make a directory (**mkdir**) and move into it (**cd**)

```
[user@cluster ~]$ pwd  
/home/user  
[user@cluster ~]$ mkdir training  
[user@cluster ~]$ cd training/
```

processes

You can see your current processes using **ps**.

```
[user@cluster training]$ ps
  PID TTY          TIME CMD
 14859 pts/22    00:00:00 bash
 18511 pts/22    00:00:00 ps
```

You can see what else *this* computer is doing using **top**

```
[user@cluster training]$ top
```

top output

top uses the whole screen. Type 'q' to get your screen back.

```
top - 16:26:38 up 58 days, 22:33, 36 users,  load average: 0.12, 0.14, 0.12
Tasks: 618 total,   1 running, 617 sleeping,   0 stopped,   0 zombie
Cpu(s):  0.1%us,  0.2%sy,  0.0%ni, 99.5%id,  0.2%wa,  0.0%hi,  0.0%si,  0.0%st
Mem:  16437908k total, 10473016k used,  5964892k free,  2611564k buffers
Swap: 16779852k total,  162896k used, 16616956k free,  2158536k cached
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
975	root	0	-20	22712	3832	2196	S	1	0.0	28:44.67	lim
4686	root	15	0	0	0	0	S	0	0.0	3:11.36	nfsd
19175	user	15	0	11048	1592	864	R	0	0.0	0:00.14	top
1	root	15	0	10364	600	564	S	0	0.0	0:12.04	init
2	root	RT	-5	0	0	0	S	0	0.0	0:04.62	migration/0
3	root	34	19	0	0	0	S	0	0.0	0:00.41	ksoftirqd/0
4	root	RT	-5	0	0	0	S	0	0.0	0:00.00	watchdog/0
5	root	RT	-5	0	0	0	S	0	0.0	0:07.06	migration/1
6	root	34	19	0	0	0	S	0	0.0	0:00.81	ksoftirqd/1
7	root	RT	-5	0	0	0	S	0	0.0	0:00.00	watchdog/1
8	root	RT	-5	0	0	0	S	0	0.0	0:05.66	migration/2
9	root	34	19	0	0	0	S	0	0.0	0:01.43	ksoftirqd/2
10	root	RT	-5	0	0	0	S	0	0.0	0:00.00	watchdog/2
11	root	RT	-5	0	0	0	S	0	0.0	0:05.09	migration/3
12	root	34	19	0	0	0	S	0	0.0	0:00.73	ksoftirqd/3
13	root	RT	-5	0	0	0	S	0	0.0	0:00.00	watchdog/3
14	root	RT	-5	0	0	0	S	0	0.0	0:20.33	migration/4

The 'sleep' command

The **sleep** command doesn't do much – but you can control how many seconds it does it for, and it doesn't use much CPU or I/O

```
[user@cluster training]$ sleep 10  
[user@cluster training]$
```

Stop and suspend

If we get bored, change our mind, or think something is wrong we can interrupt jobs. To stop a job, type '^C' at the command line (that's [Ctrl]+[C] together).

```
[user@cluster training]$ sleep 100
[user@cluster training]$
```

If you don't want to stop the job, you can suspend it. Type '^Z' (that's [Ctrl]+[Z]).

```
[user@cluster training]$ sleep 100
[1]+  Stopped                  sleep 100
[user@cluster training]$
```

backgrounding

We now have a suspended job, which will never finish. To get it to carry on, we can put it in the 'background' using **bg**

```
[user@cluster training]$ sleep 100
[1]+  Stopped                  sleep 100
[user@cluster training]$ bg
[1]+  sleep 100 &
[user@cluster training]$ ps
  PID TTY          TIME CMD
 14859 pts/22    00:00:00 bash
 24799 pts/22    00:00:00 sleep
 25377 pts/22    00:00:00 ps
```

You can put a job in the background deliberately using the '&' character at the end of the command.

```
[user@cluster training]$ sleep 100 &
[1] 787
[user@cluster training]$ ps
  PID TTY          TIME CMD
   787 pts/22    00:00:00 sleep
   804 pts/22    00:00:00 ps
 14859 pts/22    00:00:00 bash
```


Killing processes

If you don't want to wait for it to finish, or think it is broken in some way, you can terminate it using the **kill** command.

Kill has a variety of gentle options to allow the process to exit gracefully. You only need to know one – signal **-9**, better known by its name **-KILL**

```
[user@cluster training]$ sleep 100 &
[1] 787
[user@cluster training]$ ps
  PID TTY          TIME CMD
   787 pts/22    00:00:00 sleep
   804 pts/22    00:00:00 ps
 14859 pts/22    00:00:00 bash
[user@cluster training]$ kill -KILL 787
[user@cluster training]$
[1]+  Killed                  sleep 100
[user@cluster training]$
```

A simple example

Sleep is a good example, but it doesn't produce any output. We want to wrap it up with messages – in unix you use **echo** to do this.

The colon here allows us to put multiple commands on a single line.

```
[user@cluster training]$ echo start; sleep 1; echo finish  
start  
finish  
[user@cluster training]$
```

Creating a script

Cluster programming makes use of scripts, so we'll turn this list of commands into a script.

You can type directly into a file using **cat** if you know that the end of file character is a '^D'.

```
[user@cluster training]$ cat > script.sh
echo start
sleep 10
echo finish
[user@cluster training]$
```

You can run a script by executing **bash** **<scriptname>** or by making it directly executable with **chmod**. The **./** is important – the shell only looks for executables in certain places – the **'PATH'**.

```
[user@cluster training]$ chmod +x script.sh
[user@cluster training]$ ./script.sh
start
finish
```

Running the script

Now we are ready to start running our script, or sending it as a cluster job.

```
[user@cluster training]$ ./script.sh > script.out &
[1] 7594
[user@cluster training]$ ps
  PID TTY          TIME CMD
 7594 pts/22    00:00:00 bash
 7595 pts/22    00:00:00 sleep
 7598 pts/22    00:00:00 ps
14859 pts/22    00:00:00 bash
[user@cluster training]$
[1]+  Done                  ./script.sh > script.out
```



CANCER
RESEARCH
UK

CAMBRIDGE
INSTITUTE

Cluster practical sessions

Part II – cluster job submission



UNIVERSITY OF
CAMBRIDGE

Submitting a job

Now we know enough to run our script on the cluster.

Simply submit the job using **bsub**.

Notice

- the output file, in a **/lustre** directory - **/home** isn't writeable from cluster nodes.
- The script is fully located (and recall we made it executable). Each job gets a fresh shell, initially located in **/home**.

```
[user@cluster training]$ bsub -o /lustre/computing/user/training/script.out  
'~/training/script.sh'  
No -Rrusage[mem=<memoryreservation in MB>] reservation request has been made  
setting rusage and -M to default values of 2GB  
Job <474391> is submitted to default queue <cluster>.  
[user@cluster training]$
```

Look at running jobs

While the job is running, you can see it with **bjobs**.

```
[user@cluster training]$ bjobs
JOBID  USER   STAT  QUEUE   FROM_HOST   EXEC_HOST   JOB_NAME   SUBMIT_TIME
474391  user    RUN   cluster uk-cri-lcst crinode52   *script.sh May 14 16:55
[user@cluster training]$
```

Once it's finished, you can see the output.

```
[user@cluster training]$ bjobs
No unfinished job found
[user@cluster training]$ cat /lustre/computing/user/training/script.out
```

```
[user@cluster training]$ cat /lustre/computing/user/training/script.out
Sender: LSF System <lsfadmin@crinode52>
Subject: Job 474391: <~/training/script.sh> Done
```

```
Job <~/training/script.sh> was submitted from host <uk-cri-lcst01> by
user <user> in cluster <uk-cri-cluster01>.
Job was executed on host(s) <crinode52>, in queue <cluster>, as user
<user> in cluster <uk-cri-cluster01>.
</home/user> was used as the home directory.
</home/user/training> was used as the working directory.
Started at Thu May 14 16:55:10 2015
Results reported at Thu May 14 16:55:20 2015
```

Your job looked like:

```
-----
# LSBATCH: User input
~/training/script.sh
-----
```

Successfully completed.

Resource usage summary:

CPU time	:	0.03 sec.
Max Memory	:	2 MB
Max Swap	:	27 MB
Max Processes	:	1
Max Threads	:	1

The output (if any) follows:

```
start
finish
```


An alternative way to submit

Read the commands from **stdin** using '<'

```
[user@cluster training]$ bsub -o /lustre/computing/user/training/script.out <
script.sh
```

How is this different?

```
[user@cluster training]$ cat /lustre/computing/user/training/script.out
Sender: LSF System <lsfadmin@crinode98>
Subject: Job 474618: <echo start;sleep 10;echo finish> Done
[...]
Your job looked like:

-----
# LSBATCH: User input
echo start
sleep 10
echo finish

-----

[...]
The output (if any) follows:

start
finish
```

Killing a job

Just as for processes, but using **bkill**

```
[user@cluster training]$ bsub -o /lustre/computing/user/training/script.out <
script.sh
No -Rrusage[mem=<memoryreservation in MB>] reservation request has been made
setting rusage and -M to default values of 2GB
Job <474719> is submitted to default queue <cluster>.
[user@cluster training]$ bkill -s KILL 474719
Job <474719> is being terminated
[user@cluster training]$
[user@cluster training]$ bjobs
No unfinished job found
[user@cluster training]$
```

Killing isn't bad...

The scheduler manages the shutdown and still records details of the job.

```
[user@cluster training]$ cat /lustre/computing/user/training/script.out
```

```
[...]
```

```
Your job looked like:
```

```
-----  
# LSBATCH: User input  
echo start  
sleep 10  
echo finish  
-----
```

```
TERM_OWNER: job killed by owner.  
Exited with exit code 130.
```

```
Resource usage summary:
```

CPU time	:	0.02 sec.
Max Memory	:	2 MB
Max Swap	:	27 MB
Max Processes	:	1
Max Threads	:	1

```
The output (if any) follows:
```

```
start
```

Basic parallelism

Now we're ready to use the cluster at full power!

The simplest way to do this is simply running multiple jobs at once. Here we use the **bash** **'for...; do ...; done'** construct, the **bsub -J** option and the **%J** environment variable.

```
[user@cluster training]$ for j in 1 2 3; do bsub -o
/lustre/computing/user/training/script-%J.out -J job$j '~/training/script.sh';
done
No -Rusage[mem=<memoryreservation in MB>] reservation request has been made
setting rusage and -M to default values of 2GB
Job <474824> is submitted to default queue <cluster>.
No -Rusage[mem=<memoryreservation in MB>] reservation request has been made
setting rusage and -M to default values of 2GB
Job <474825> is submitted to default queue <cluster>.
No -Rusage[mem=<memoryreservation in MB>] reservation request has been made
setting rusage and -M to default values of 2GB
Job <474826> is submitted to default queue <cluster>.
[user@cluster training]$
[user@cluster training]$ bjobs
No unfinished job found
[user@cluster training]$ ls /lustre/computing/user/training/
script-474824.out  script-474825.out  script-474826.out
```



CANCER
RESEARCH
UK

CAMBRIDGE
INSTITUTE

Fin



UNIVERSITY OF
CAMBRIDGE