



CANCER
RESEARCH
UK

CAMBRIDGE
INSTITUTE

CRUK cluster introduction (II of III)

Using the scheduler for job
submission



UNIVERSITY OF
CAMBRIDGE

Overview

Our job scheduling system is Platform LSF

In their own words:

'LSF provides the capabilities to manage and accelerate workload processing across heterogeneous distributed compute environments. It is comprised of a comprehensive set of intelligent scheduling policies to ensure that the right resources are automatically allocated to the right jobs, for maximum application performance and efficiency.'

Version 7.0.4

LSF will allow you to:

- Submit jobs to the cluster.
- Specify which queue to submit your jobs to.
- Request memory resources for your jobs.
- Set memory limits for your jobs.
- Check the status of the jobs you have submitted.
- Check the status of the hosts within the cluster.
- Kill jobs that you have submitted to the cluster.

The most useful LSF commands

These are:

- **bsub**
- **bjobs**
- **bhosts**
- **bqueues**
- **bkill**

Type command followed by -h for usage details.

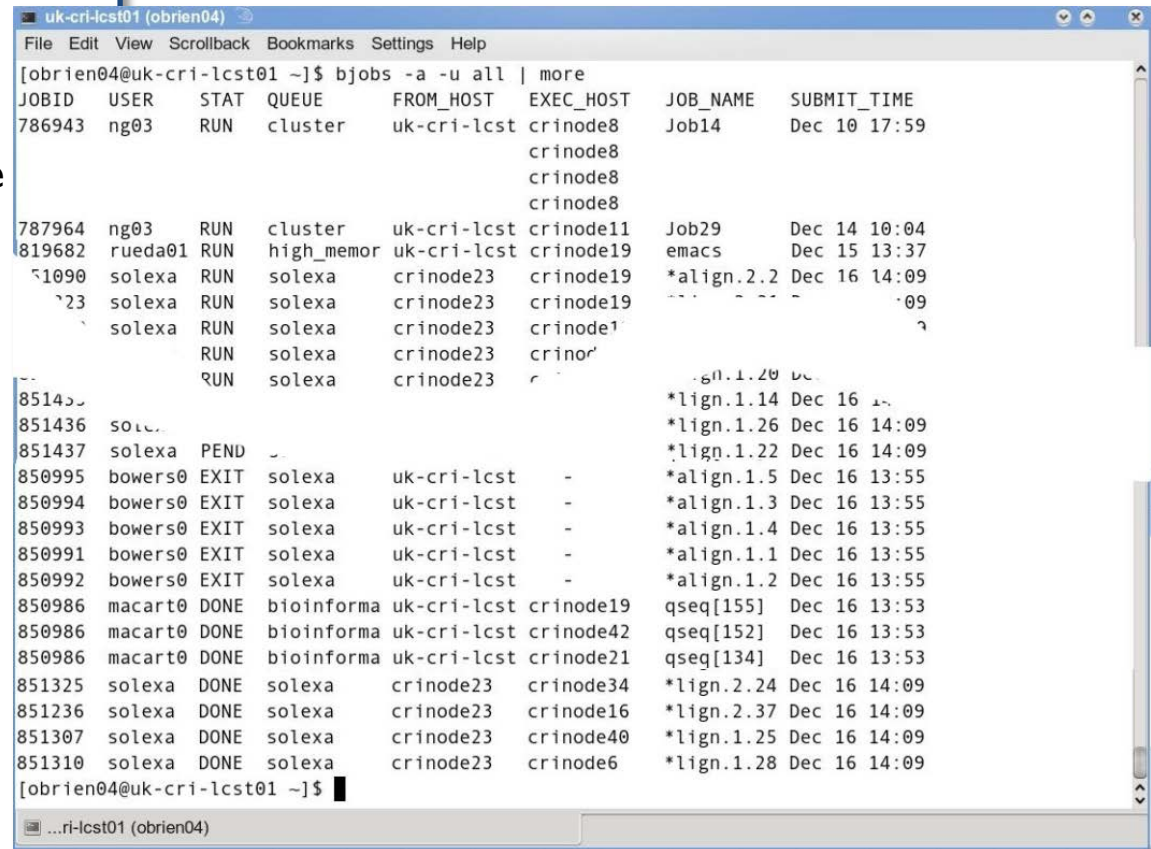
Many LSF commands have shared syntax, i.e. the **-l** (long output), **-u** (User) and **-m** (host) flags.

A Job

A command or series of commands submitted to the cluster with associated resource requirements and limits.

Status of jobs running on the cluster can be seen with the command

uk-cri-lcst01 > bjobs -a -u all



```
uk-cri-lcst01 (obrien04)
File Edit View Scrollback Bookmarks Settings Help
[obrien04@uk-cri-lcst01 ~]$ bjobs -a -u all | more
JOBID  USER  STAT  QUEUE      FROM_HOST  EXEC_HOST  JOB_NAME    SUBMIT_TIME
786943  ng03    RUN   cluster    uk-cri-lcst  crinode8   Job14       Dec 10 17:59
                                     crinode8
                                     crinode8
787964  ng03    RUN   cluster    uk-cri-lcst  crinode11  Job29       Dec 14 10:04
819682  rueda01 RUN   high_memor uk-cri-lcst  crinode19  emacs       Dec 15 13:37
51090   solexa  RUN   solexa     crinode23   crinode19  *align.2.2  Dec 16 14:09
123     solexa  RUN   solexa     crinode23   crinode19  ...         09
       solexa  RUN   solexa     crinode23   crinode1
       solexa  RUN   solexa     crinode23   crino
       solexa  RUN   solexa     crinode23   crino
851435  solexa  RUN   solexa     crinode23   crino
851436  solexa  RUN   solexa     crinode23   crino
851437  solexa  PEND  -          -          -          *lign.1.26 Dec 16 14:09
850995  bowers0 EXIT  solexa     uk-cri-lcst  -          *align.1.5 Dec 16 13:55
850994  bowers0 EXIT  solexa     uk-cri-lcst  -          *align.1.3 Dec 16 13:55
850993  bowers0 EXIT  solexa     uk-cri-lcst  -          *align.1.4 Dec 16 13:55
850991  bowers0 EXIT  solexa     uk-cri-lcst  -          *align.1.1 Dec 16 13:55
850992  bowers0 EXIT  solexa     uk-cri-lcst  -          *align.1.2 Dec 16 13:55
850986  macart0 DONE  bioinforma uk-cri-lcst  crinode19  qseq[155]  Dec 16 13:53
850986  macart0 DONE  bioinforma uk-cri-lcst  crinode42  qseq[152]  Dec 16 13:53
850986  macart0 DONE  bioinforma uk-cri-lcst  crinode21  qseq[134]  Dec 16 13:53
851325  solexa  DONE  solexa     crinode23   crinode34  *lign.2.24 Dec 16 14:09
851236  solexa  DONE  solexa     crinode23   crinode16  *lign.2.37 Dec 16 14:09
851307  solexa  DONE  solexa     crinode23   crinode40  *lign.1.25 Dec 16 14:09
851310  solexa  DONE  solexa     crinode23   crinode6   *lign.1.28 Dec 16 14:09
[obrien04@uk-cri-lcst01 ~]$
```

The bsub command

A monster, even by unix standards.

```
bsub [ -h ] [ -V ] [ -H ] [ -x ] [ -r ] [ -N ] [ -B ]  
[ -I | -Ip | -Is | -K ]  
[ -T time_event ]  
[ [ -X "exception_cond([params])::action" ] ... ]  
[ -w depend_cond ]  
[ -q queue_name ... ] [ -a application_name ]  
[ -m host_name[+[pref_level]] ... ]  
[ -n min_proc[,max_proc] ]  
[ -R res_req ]  
[ -J job_name_spec ] [ -b begin_time ] [ -t term_time  
]  
[ -i in_file ] [ -o out_file ] [ -e err_file ]  
[ -u mail_user ] [ [ -f "lfile op [ rfile ]" ] ... ]  
[ -E "pre_exec_command [ argument ... ]" ]  
[ -c cpu_limit[/host_spec] ] [ -W  
run_limit[/host_spec] ]  
[ -F file_limit ] [ -M mem_limit ] [ -D data_limit ]  
[ -S stack_limit ] [ -C core_limit ]  
[ -k "chkpnt_dir [ chkpnt_period ]" ] [ -w  
depend_cond ]  
[ -L login_shell ] [ -P project_name ]  
[ -G user_group ] [ command [ argument ... ] ]
```

A Queue:

A queue for job submissions associated with specified users and cluster hosts, and providing specified default resources.

You specify the queue to use by adding the **-q <queuenam>** option to the bsub command.

Most CRI cluster queues have a default memory resource limit of 2GB per job. You can override the default by adding the **-M <memory size in KB>** to the bsub command.

An example submission request to the queue named cluster to run a job which overrides the 2GB memory limit allowing the job to use 4GBs and requesting that the host has/hosts have 4GB memory available, may look like this:

```
uk-cri-lcst01 > bsub -q cluster -M 4194304  
-R "rusage[mem=4096]" <command>
```

```
uk-cri-lcst01 (obrien04)
File Edit View Scrollback Bookmarks Settings Help
[obrien04@uk-cri-lcst01 ~]$ bqueues
QUEUE_NAME      PRIO STATUS      MAX JL/U JL/P JL/H NJOBS  PEND  RUN  SUSP
cluster         30  Open:Active   240  -   -   -    77    0    77    0
solexa          30  Open:Active   240  -   -   -   430   273  157    0
genomics        30  Open:Active   240  -   -   -    0     0    0     0
mri             30  Open:Active   240  -   -   -    0     0    0     0
stlab           30  Open:Active   240  -   -   -    0     0    0     0
bioinformatics  30  Open:Active   240  -   -   -   12     3    9     0
information_sys 30  Open:Active   240  -   -   -    0     0    0     0
high_memory     30  Open:Active    -   -   -   -    1     0    1     0
groundfloor     4   Open:Active    2   -   -   -    0     0    0     0
basement        3   Open:Active    2   -   -   -    0     0    0     0
test            2   Open:Active    8   -   -   -    0     0    0     0
[obrien04@uk-cri-lcst01 ~]$
```

...ri-lcst01 (obrien04)

The 'test' queue

We also have a test queue, called test, which comprises one sacrificial host (8 cores).

Please submit brand new jobs, and jobs that you feel might have become rogue after tweaking, to this queue before setting them loose on the cluster. If it kills the host, let us know and we'll restart it ready for further testing.

(Of course, if your job does kill the test host, please don't release it to the cluster in general until it has been debugged and shown to behave itself in the test environment.)

Viewing the test queue properties

Use the long version of the bqueues command.

uk-cri-lcst01 > bqueues -l test

```
uk-cri-lcst01 (obrien04)
File Edit View Scrollback Bookmarks Settings Help
[obrien04@uk-cri-lcst01 ~]$ bqueues -l test

QUEUE: test
  -- Test queue with single sacrificial host, for debugging new jobs.

PARAMETERS/STATISTICS
PRIO NICE STATUS      MAX JL/U JL/P JL/H NJOBS  PEND  RUN  SSUSP  USUSP  RSV
  2   20  Open:Active      8   -   -   -     0     0     0     0     0     0

DEFAULT LIMITS:
MEMLIMIT
2097152 K

MAXIMUM LIMITS:
MEMLIMIT
16777216 K

SCHEDULING PARAMETERS
          r15s  r1m  r15m  ut      pg    io    ls    it    tmp    swp    mem
loadSched -    -    -    -      -    -    -    -    -    -    -    -
loadStop  -    -    -    -      -    -    -    -    -    -    -    -

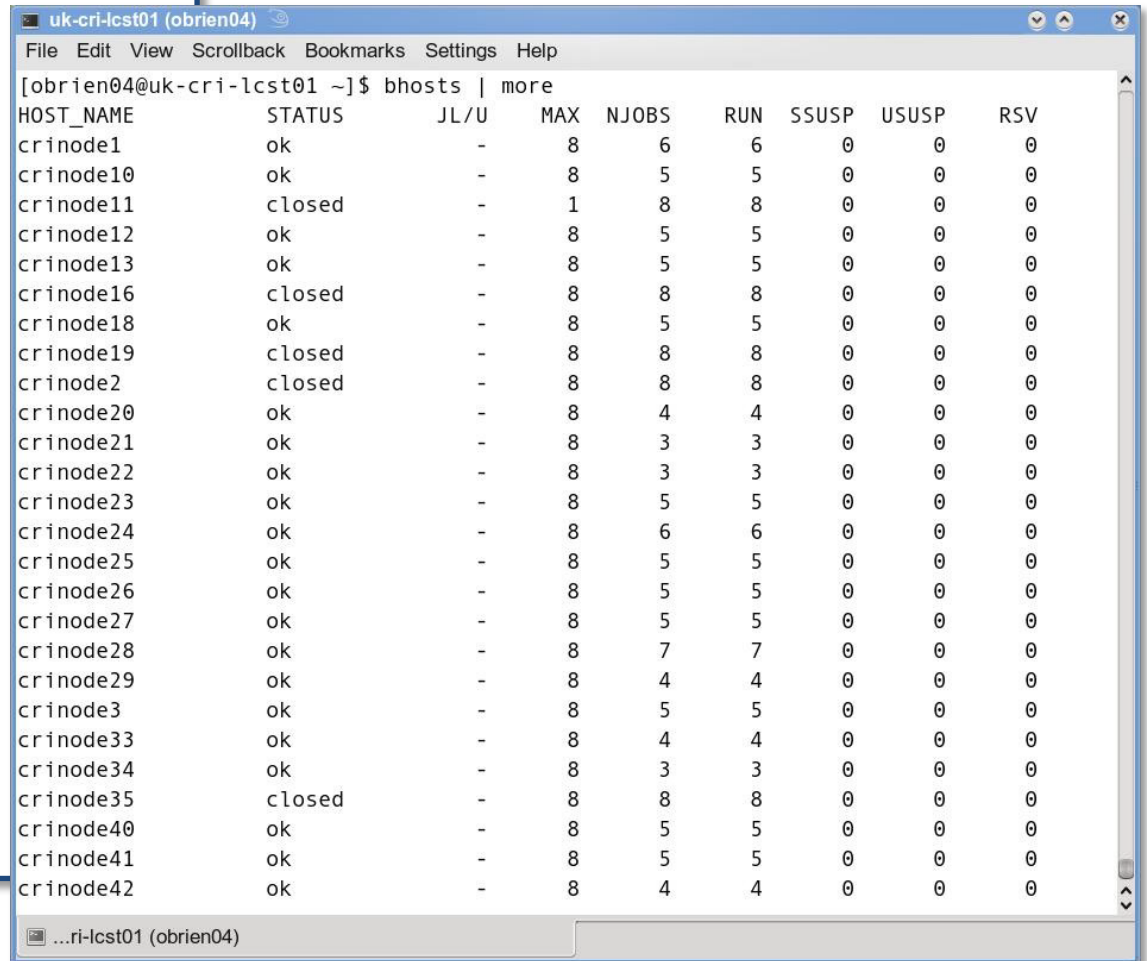
USERS: all
HOSTS: test/
PREEMPTION: PREEMPTABLE

[obrien04@uk-cri-lcst01 ~]$
```

Checking the status of hosts within the cluster

Check hosts status using the bhosts command

uk-cri-lcst01 > bhosts



HOST_NAME	STATUS	JL/U	MAX	NJOBS	RUN	SSUSP	USUSP	RSV
crinode1	ok	-	8	6	6	0	0	0
crinode10	ok	-	8	5	5	0	0	0
crinode11	closed	-	1	8	8	0	0	0
crinode12	ok	-	8	5	5	0	0	0
crinode13	ok	-	8	5	5	0	0	0
crinode16	closed	-	8	8	8	0	0	0
crinode18	ok	-	8	5	5	0	0	0
crinode19	closed	-	8	8	8	0	0	0
crinode2	closed	-	8	8	8	0	0	0
crinode20	ok	-	8	4	4	0	0	0
crinode21	ok	-	8	3	3	0	0	0
crinode22	ok	-	8	3	3	0	0	0
crinode23	ok	-	8	5	5	0	0	0
crinode24	ok	-	8	6	6	0	0	0
crinode25	ok	-	8	5	5	0	0	0
crinode26	ok	-	8	5	5	0	0	0
crinode27	ok	-	8	5	5	0	0	0
crinode28	ok	-	8	7	7	0	0	0
crinode29	ok	-	8	4	4	0	0	0
crinode3	ok	-	8	5	5	0	0	0
crinode33	ok	-	8	4	4	0	0	0
crinode34	ok	-	8	3	3	0	0	0
crinode35	closed	-	8	8	8	0	0	0
crinode40	ok	-	8	5	5	0	0	0
crinode41	ok	-	8	5	5	0	0	0
crinode42	ok	-	8	4	4	0	0	0

check host resources

Use the lsload command

uk-cri-lcst01 > lsload -R mem

```
uk-cri-lcst01 (obrien04)
File Edit View Scrollback Bookmarks Settings Help
[obrien04@uk-cri-lcst01 ~]$ lsload -R mem
HOST_NAME      status  r15s  r1m  r15m  ut    pg  ls    it    tmp  swp  mem
crinode61      ok      0.0   0.1   0.1   0%    0.0  1 13200  71G  16G  15G
crinode54      ok      0.0   0.0   0.0   0%    0.2  0 44896  71G  12G  11G
crinode7       ok      0.0   0.0   0.0   0%    0.3  0 18816  71G  13G  10G
crinode35      ok      0.0   0.2   0.0   0%    0.4  0 44800  71G  13G  11G
crinode2       ok      0.0   0.0   0.0   0%    0.2  1 31840  71G
lcst01         ok      0.0   0.0   1.5   14%   0.0  10    0   71G
crinode45      ok      6.3   6.7   5.4   73%   0.0  0 54784  71G  16G  5420M
crinode20      ok      6.3   5.8   5.0   73%   0.0  0 54784  71G  16G  7000M
crinode19      ok      6.5   7.9   8.1   96%   0.0  0 31968  71G  16G  16G
crinode52      ok      6.7   7.0   7.8   87%   0.0  0 11664  71G  16G  18G
crinode16      ok      6.9   7.2   7.5   87%   0.0  0 19056  71G  16G  13G
crinode59      ok      7.0   7.0   5.4   88%   0.0  0 2e+05  71G  16G  3676M
crinode48      ok      8.9   7.9   7.8   98%   0.0  0 44896  71G  16G  10G
[obrien04@uk-cri-lcst01 ~]$
```

Killing jobs with bkill

You can kill your own jobs if they appear not to be running as intended.

LSF also allows for jobs to be stopped and restarted (provided they were submitted with the **bsub -r** flag).

```
bstop [ -h ] [ -V ] [ -q queue_name ] [ -m host_name ]  
[ -u user_name | all ] [ -J job_name ] [ jobId |  
"jobId[index_list]" ... ]  
  
bresume [ -h ] [ -V ] [ -q queue_name ] [ -m host_name ]  
[ -u user_name | all ] [ -J job_name ] [ jobId |  
"jobId[index_list]" ... ]  
  
bkill [ -h ] [ -V ] [ -l ] [ -s (signal_value | signal_name ) ]  
[ -q queue_name ] [ -m host_name ] [ -u (user_name | all) ]  
[ -J job_name ] [ jobId | "jobId[index_list]" ... ]
```

Running scripts with LSF

Script must be executable, the linux chmod command can be used to set the executable attribute

```
uk-cri-lcst01> chmod u+x  
/lustre/xxlab/xxuser/xxscript.sh
```

The script can then also be run by redirecting it to the bsub command using one of the linux redirection operators "<"

```
uk-cri-lcst01> bsub -q cluster -M 4194304  
-R "rusage[mem=4096]" < xxscript.sh
```

You can also include BSUB options within the script i.e.

```
#!/bin/sh  
#BSUB -o myoutput.log  
#BSUB -e myerror.log  
#BSUB -a R  
cd /lustre/xxlab/xxuser  
<command>
```


Simple Parallel Computing

Most HPC at CRI consists of breaking up your dataset into chunks and firing off a separate job for each chunk.

‘ Magnetisation Transfer Map Example

Here small groups of pixels from T2 weighted image slices are processed to calculate the elements of the MT map.’

(Example courtesy of Dominick McIntyre,
Griffiths Group)



Message Passing parallel jobs

Uses the generic PjL (Parallel Job Launcher) framework. You can easily recognise it because of the use of the `-a openmpi` flag and `mpirun.lsf`

```
uk-cri-lcst01 > bsub -o %J.out -e %J.err -n 4 -R "span[ptile=1]"  
-a openmpi mpirun.lsf ./test
```

In recent versions of LSF, another framework is also available, and it permits a tight (native) integration with the MPIs (this is why there is the OpenMPI integration)

```
uk-cri-lcst01 > bsub -o %J.out -e %J.err -n 4 -R "span[ptile=1]"  
mpirun ./test
```

A job submission example

from Ben Davis, bioinformatics

'This is a shell command line loop I ran recently... quite a good example as the actual thing running is just creating checksum files so doesn't distract.'

```
for f in `seq 1 16`; do bsub -n 1 -M 1048576 -R 'span[hosts=1]
select[mem>=1024 && tmp>=2000] rusage[mem=1024, tmp=2000]' -o md5cs-%J.out
-J md5cs$f -q solexa md5sum -c SJD_`$f`.md5;done;
```

'Here multiple named jobs are submitted to the solexa queue. Each checks a single file specified by numeric indices incorporated in the file title and creates a named output file where the title contains the job ID.'

Job array example

from Stewart MacArthur, bioinformatics

'Here is a self contained example of the basics of using job arrays. The main benefit in this case is the not speed but the ability to control the number of jobs running at any one time, using the %50 notation in the bsub ... I find job arrays particularly useful for running lots of small jobs, as there is only a single job submission there is little LSF overhead, compared to submitting 1000 separate jobs, which takes some time. Also being able to control the number of running jobs stops you swamping your queue with jobs and leaves space for others to get jobs running.'

```
### Generate a random big file that we want to sort, 10 Million lines
perl -e 'for (1..1E7){printf("%.0f\n",rand()*1E7)};' > bigFile
### Split the file up into chunks with 10,000 lines in each chunk
split -a 3 -d -l 10000 bigFile split
### rename the files on a 1-1000 scheme not 0-999
for f in split*;do mv ${f} $(echo ${f} |perl -ne 'm/split(0*)(\d+)/g;print
"Split", $2+1, "\n";');done
### submit a job array, allowing 50 jobs to be run at anyone time
bsub -J "sort[1-1000]%50" "sort -n Split\${LSB_JOBINDEX} >Split\
${LSB_JOBINDEX}.sorted"
### merge the sorted files together once all the jobs are finished
sort -n -m *.sorted >bigFile.sorted
### Delete the temp files
rm -f Split*
```

Job dependency

-w 'dependency_expression'

LSF does not place your job unless the dependency expression evaluates to TRUE. The dependency expression is a logical expression composed of one or more dependency conditions.

To make dependency expression of multiple conditions, use the following logical operators:
&& (AND) **||** (OR) **!** (NOT)

Use the * with dependency conditions to define one-to-one dependency among job array elements such that each element of one array depends on the corresponding element of another array. The job array size must be identical.

For example:

```
bsub      -w "done(myarrayA[*])"  
          -J "myArrayB[1-10]"  
          myJob2
```



CANCER
RESEARCH
UK

CAMBRIDGE
INSTITUTE

Practical session II



UNIVERSITY OF
CAMBRIDGE