



CANCER
RESEARCH
UK

CAMBRIDGE
INSTITUTE

CRUK cluster practical sessions (SLURM)

Part I – processes & scripts



UNIVERSITY OF
CAMBRIDGE

login

Log in to the head node, clustl-headnode, using **ssh** and your usual user name & password.

```
SSH Secure Shell 3.2.9 (Build 283)
Copyright (c) 2000-2003 SSH Communications Security Corp - http://www.ssh.com/

This copy of SSH Secure Shell is a non-commercial version.
This version does not include PKI and PKCS #11 functionality.

Last login: Mon Sep 19 10:44:07 2016 from bp7r25j.cri.camres.org
[user@cluster ~]$
```

You're ready to start.

navigate

Find out where you are using **pwd**.

Make a directory (**mkdir**) and move into it (**cd**)

```
[user@cluster ~]$ pwd  
/home/user  
[user@cluster ~]$ mkdir training  
[user@cluster ~]$ cd training/  
[user@cluster training]$
```

processes

You can see your current processes using **ps**.

```
[user@cluster training]$ ps
  PID TTY          TIME CMD
 14859 pts/22    00:00:00 bash
 18511 pts/22    00:00:00 ps
```

You can see what else *this* computer is doing using **top**

```
[user@cluster training]$ top
```

top output

top uses the whole screen. Type 'q' to get your screen back.

```
top - 16:26:38 up 58 days, 22:33, 36 users,  load average: 0.12, 0.14, 0.12
Tasks: 618 total,   1 running, 617 sleeping,   0 stopped,   0 zombie
Cpu(s):  0.1%us,   0.2%sy,   0.0%ni, 99.5%id,   0.2%wa,   0.0%hi,   0.0%si,   0.0%st
Mem:  16437908k total, 10473016k used,  5964892k free,  2611564k buffers
Swap: 16779852k total,  162896k used, 16616956k free,  2158536k cached
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
975	root	0	-20	22712	3832	2196	S	1	0.0	28:44.67	lim
4686	root	15	0	0	0	0	S	0	0.0	3:11.36	nfsd
19175	user	15	0	11048	1592	864	R	0	0.0	0:00.14	top
1	root	15	0	10364	600	564	S	0	0.0	0:12.04	init

The 'sleep' command

The **sleep** command doesn't do much – but you can control how many seconds it does it for, and it doesn't use much CPU or I/O

```
[user@cluster training]$ sleep 10  
[user@cluster training]$
```

Stop and suspend

If we get bored, change our mind, or think something is wrong we can interrupt jobs. To stop a job, type '^C' at the command line (that's [Ctrl]+[C] together).

```
[user@cluster training]$ sleep 100  
[user@cluster training]$
```

If you don't want to stop the job, you can suspend it. Type '^Z' (that's [Ctrl]+[Z]). Type 'fg' to bring the job back to the foreground.

```
[user@cluster training]$ sleep 100  
[1]+  Stopped                  sleep 100  
[user@cluster training]$ fg
```

backgrounding

When we have suspended a job (which will never finish). To get it to carry on, we can put it in the 'background' using **bg**

```
[user@cluster training]$ sleep 100
[1]+  Stopped                  sleep 100
[user@cluster training]$ bg
[1]+ sleep 100 &
[user@cluster training]$ ps
  PID TTY          TIME CMD
 14859 pts/22    00:00:00 bash
 24799 pts/22    00:00:00 sleep
 25377 pts/22    00:00:00 ps
```

You can put a job in the background deliberately using the '&' character at the end of the command.

```
[user@cluster training]$ sleep 100 &
[1] 787
[user@cluster training]$ ps
  PID TTY          TIME CMD
   787 pts/22    00:00:00 sleep
   804 pts/22    00:00:00 ps
 14859 pts/22    00:00:00 bash
```


Killing processes

If you don't want to wait for it to finish, or think it is broken in some way, you can terminate it using the **kill** command.

Kill has a variety of gentle options to allow the process to exit gracefully. If these fail one – signal **-9**, or **-KILL** will normally remove the process.

```
[user@cluster training]$ sleep 100 &
[1] 787
[user@cluster training]$ ps
  PID TTY          TIME CMD
   787 pts/22    00:00:00 sleep
   804 pts/22    00:00:00 ps
 14859 pts/22    00:00:00 bash
[user@cluster training]$ kill -KILL 787
[user@cluster training]$
[1]+  Killed                  sleep 100
[user@cluster training]$
```

A simple example

Sleep is a good example, but it doesn't produce any output. We want to wrap it up with messages – in unix you use **echo** to do this.

The colon here allows us to put multiple commands on a single line.

```
[user@cluster training]$ echo start; sleep 1; echo finish
start
finish
[user@cluster training]$
```

Creating a script

Cluster programming makes use of scripts, so we'll turn this list of commands into a script.

You can type directly into a file using **cat** if you know that the end of file character is a '^D'.

```
[user@cluster training]$ cat > script.sh
#!/usr/bin/bash
echo start
sleep 10
echo finish
[user@cluster training]$
```

You can run a script by executing **bash** **<scriptname>** or by making it directly executable with **chmod**. The **./** is important – the shell only looks for executables in certain places – the **'PATH'**.

```
[user@cluster training]$ chmod u+x script.sh
[user@cluster training]$ ./script.sh
start
finish
```

Running the script

Now we are ready to start running our script, or sending it as a cluster job.

```
[user@cluster training]$ ./script.sh > script.out &
[1] 7594
[user@cluster training]$ ps
  PID TTY          TIME CMD
 7594 pts/22    00:00:00 bash
 7595 pts/22    00:00:00 sleep
 7598 pts/22    00:00:00 ps
14859 pts/22    00:00:00 bash
[user@cluster training]$
[1]+  Done                  ./script.sh > script.out
```



CANCER
RESEARCH
UK

CAMBRIDGE
INSTITUTE

Cluster practical sessions

Part II – cluster job submission



UNIVERSITY OF
CAMBRIDGE

Submitting a job

Now we know enough to run our script on the cluster.

Simply submit the job using **sbatch**.

- the output file is written to a **Lustre file system** directory
- Create directory with username if it doesn't exist – e.g. `mkdir /scratcha/stlab/garret01`
- **/home** is writeable from cluster nodes, but won't perform as well.

```
[user@cluster training]$ sbatch --output=/scratcha/stlab/garret01/%N-%j.out script.sh
Submitted batch job 200875
```

- All read and write operations from within jobs running on nodes should use either `/scratchb` or `/scratcha` directories.

Look at running jobs

While the job is running, you can see it with **squeue**.

```
[user@cluster training]$ squeue
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST (REASON)
200876	general	script.s	user	R	0:02	1	clust1-node-3
200867	general	bash	sawle01	R	4:01:05	1	clust1-node-2
175393	general	MB99.6.v	eldrid01	R	9-00:28:46	1	clust1-node-30
175330	general	vardict_	eldrid01	R	9-01:40:03	1	clust1-node-1

Once it's finished, you can see the output.

```
[user@cluster training]$ squeue
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST (REASON)
200867	general	bash	sawle01	R	4:01:33	1	clust1-node-2
175393	general	MB99.6.v	eldrid01	R	9-00:29:14	1	clust1-node-30
175330	general	vardict_	eldrid01	R	9-01:40:31	1	clust1-node-1

```
[user@cluster training]$ ls /scratch/group/user/  
clust1-node-3-200877.out
```

What happened?

The output went into the file as expected:

```
[user@cluster training]$ cat /scratcha/group/user/clust1-node-3-200877.out
start
finish
```

Other information is stored, and available via sacct:

```
[user@cluster training]$ sacct -j 200877
```

JobID	JobName	Partition	Account	AllocCPUS	State	ExitCode
200877	script.sh	general	group	1	COMPLETED	0:0
200877.batch	batch		group	1	COMPLETED	0:0


```
[user@cluster training]$ sacct -j 200877 --format JobID,MaxRSS,State,AllocCPUS
```

JobID	MaxRSS	State	AllocCPUS
200877		COMPLETED	1
200877.batch	2012K	COMPLETED	1

An alternative way to submit

You can submit a job directly to SLURM with **srun**. This still requires resources – it's more commonly used as part of an existing job.

```
[user@cluster training]$ srun /usr/bin/bash script.sh  
start  
finish
```

You can also use this to generate an interactive session:

```
[user@cluster training]$ srun --pty /usr/bin/bash  
[user@clust1-node-3 training]$
```

Killing a job

Just as for processes, but using **scancel**

```
[user@cluster training]$ sbatch --output=/scratcha/group/user/%N-%j.out script.sh
Submitted batch job 200889
[user@cluster training]$ squeue
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST (REASON)
200889	general	script.s	user	R	0:02	1	clust1-node-3
200867	general	bash	sawle01	R	4:28:21	1	clust1-node-2
175393	general	MB99.6.v	eldrid01	R	9-00:56:02	1	clust1-node-30
175330	general	vardict_	eldrid01	R	9-02:07:19	1	clust1-node-1

```
[user@cluster training]$ scancel 200889
[user@cluster training]$ squeue
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST (REASON)
200867	general	bash	sawle01	R	4:28:21	1	clust1-node-2
175393	general	MB99.6.v	eldrid01	R	9-00:56:02	1	clust1-node-30
175330	general	vardict_	eldrid01	R	9-02:07:19	1	clust1-node-1

NOTE: Do not use **skill** it is **NOT** a SLURM command!

Killing isn't bad...

The scheduler manages the shutdown and still records details of the job.

```
[user@cluster training]$ sacct -j 200889
```

JobID	JobName	Partition	Account	AllocCPUS	State	ExitCode
200889	script.sh	general	group	1	CANCELLED+	0:0
200889.batch	batch		group	1	CANCELLED	0:15

Basic parallelism

Now we're ready to use the cluster at full power!

One way to do this is with a job array. You can create one of these using the

--array=1-N syntax in **sbatch**

```
[user@cluster training]$ sbatch --array=1-10 --output=/scratcha/group/user/%N-%j.out
script.sh
Submitted batch job 200900
[user@cluster training]$ ls / scratcha/group/user
clust1-node-10-200908.out  clust1-node-12-200900.out  clust1-node-4-200902.out
clust1-node-10-200904.out  clust1-node-12-200906.out  clust1-node-11-200909.out
clust1-node-3-200901.out  clust1-node-5-200903.out  clust1-node-7-200905.out
clust1-node-9-200907.out
```

Or using the **srun** with the **-n** or **-N** parameters.

```
[user@cluster training]$ srun -n srun -n 41 hostname
clust1-node-9.cri.camres.org
...
clust1-node-13.cri.camres.org
[user@cluster training]$
[user@cluster training]$ srun -N 3 hostname
clust1-node-19.cri.camres.org
clust1-node-25.cri.camres.org
clust1-node-8.cri.camres.org
[user@cluster training]$
```



CANCER
RESEARCH
UK

CAMBRIDGE
INSTITUTE

Ende



UNIVERSITY OF
CAMBRIDGE